A collection of various toy cars and trucks, including a pink sedan, a yellow taxi with 'ADAC - STRASSEN WACHT' text, a white car with red racing stripes, a blue sedan, a red car, a white van, and a yellow construction vehicle, all arranged on a dark, textured surface.

Stanford Cars Dataset

NOA CHU ARTZT

ÍNDICE

Dataset

Preprocesado

Modelos

LeNet

AlexNet

VGG16

GoogleNet y Inception

ResNet

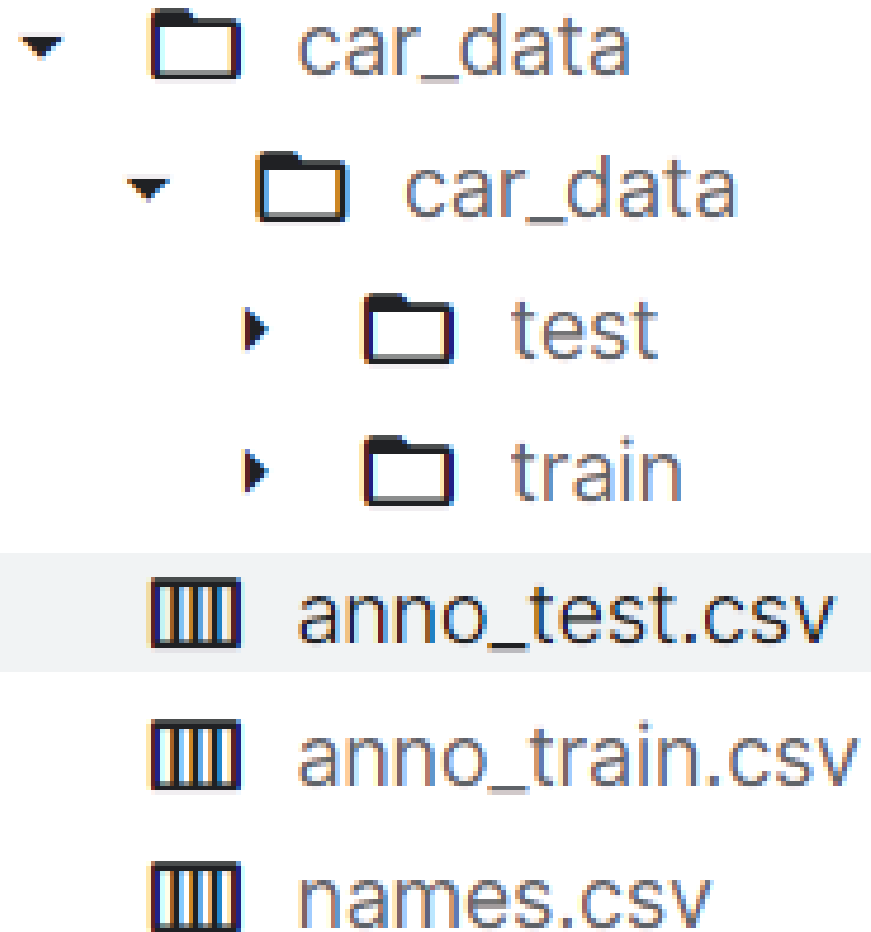
Autoencoder



Kaggle - Stanford Car Dataset by classes folder

Kaggle - Stanford Car Dataset by classes folder

- 16.185 imágenes
- 196 clases
 - 8.144 de entrenamiento
 - 8.041 de testeo





Preprocesado

```
TARGET_SIZE = 227, 224, 64, 32
```

```
BATCH_SIZE = 128, 64, 32, 16, 8
```

```
ImageDataGenerator(rescale, rotation_range,  
width_shift_range, height_shift_range, horizontal_flip,  
zoom_range, fill_mode, validation_split)
```

```
flow_from_directory(train_path, target_size, class_mode,  
subset)
```




Modelos CNN

LENET

Loss: 5.06

Accuracy: 2.98%



LENET

ACTUAL: Audi A5 Coupe 2012



ACTUAL: Dodge Challenger SRT8 2011




ACTUAL: Hyundai Tucson SUV 2012



PREDICTION: Lamborghini Gallardo LP 570-4 Superleggera 2012



PREDICTION: Nissan Leaf Hatchback 2012

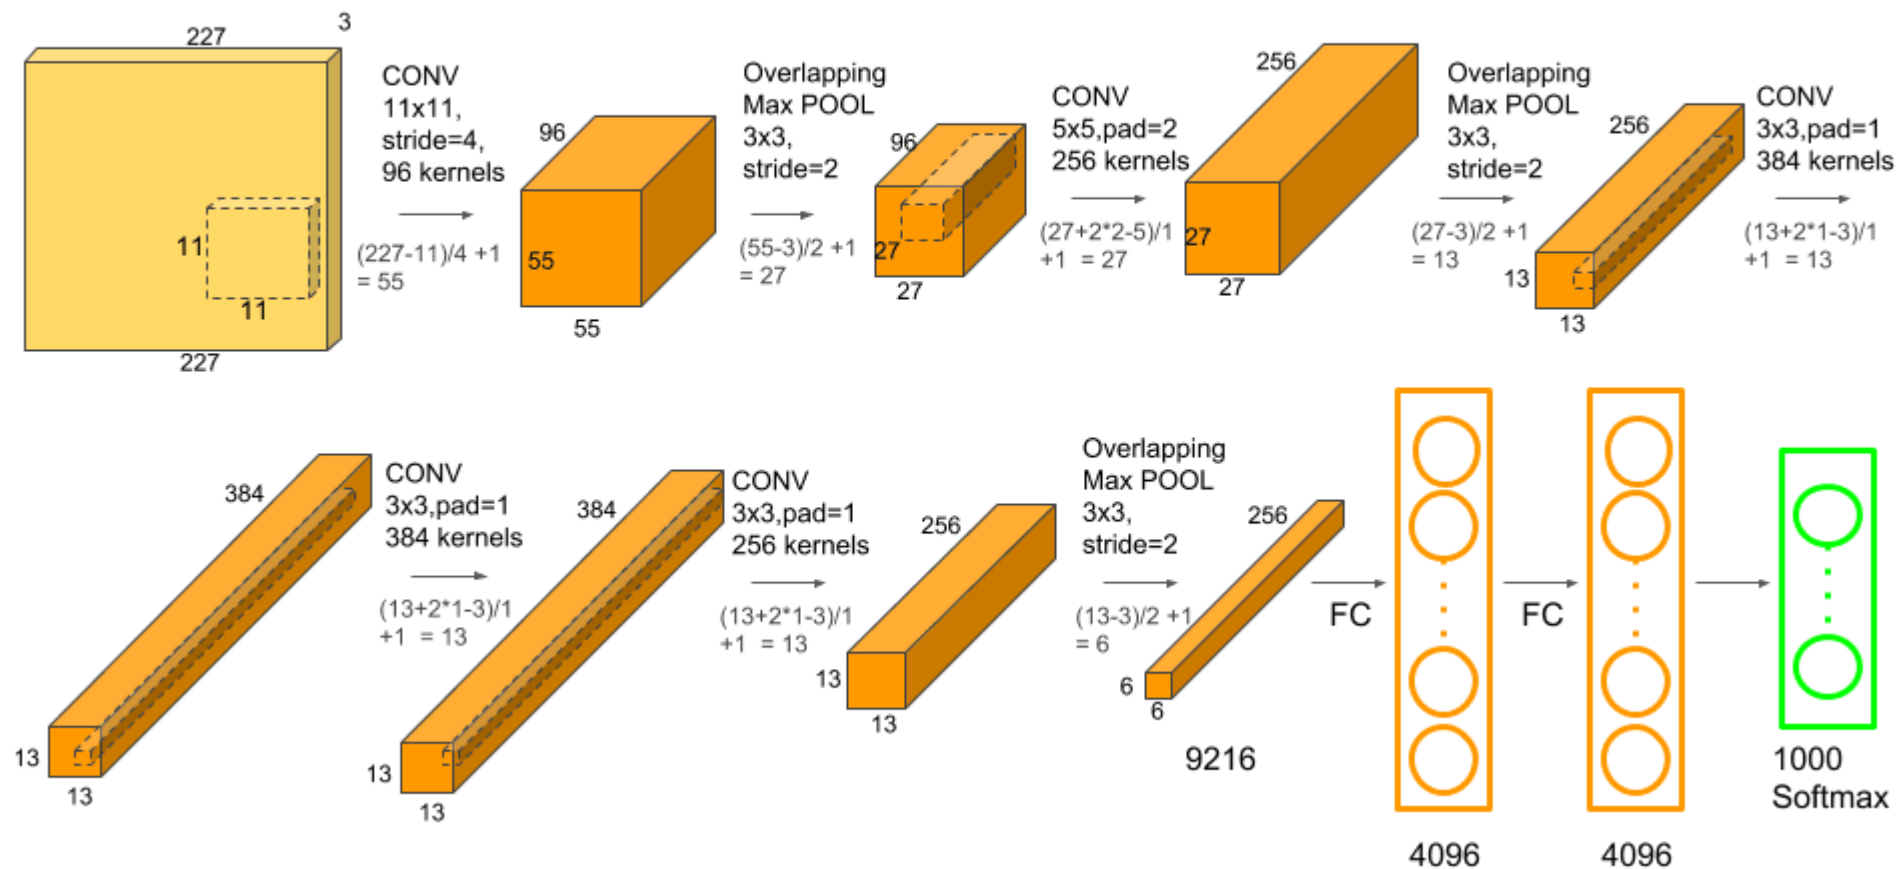


PREDICTION: Audi TT Hatchback 2011



ALEXNET

Loss: 5.27
Accuracy: 0.83%



ALEXNET

ACTUAL: Ford Edge SUV 2012



ACTUAL: Ferrari 458 Italia Coupe 2012



ACTUAL: smart fortwo Convertible 2012



PREDICTION: GMC Yukon Hybrid SUV 2012



PREDICTION: GMC Yukon Hybrid SUV 2012



PREDICTION: GMC Yukon Hybrid SUV 2012



RESNET

Loss: 5.28

Accuracy: 0.85%

GoogleNet + Inception



Loss: 5.28

Accuracy: 0.85%

VGG-16

Loss: 5.27

Accuracy: 0.83%



Transfer Learning

RESNET-50

Loss: 1.12

Accuracy: 70.65%

```
# Load pre-trained ResNet50 model without top layers
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=self.input_shape)

# Freeze base model layers
for layer in base_model.layers:
    if "BatchNormalization" in layer.__class__.__name__:
        layer.trainable = True

# Add custom layers for classification
x = GlobalAveragePooling2D()(base_model.output)
x = BatchNormalization()(x)

x = Dense(1024, activation='relu')(x)
x = Dropout(0.5)(x)
x = BatchNormalization()(x)

x = Dense(512, activation='relu')(x)
x = Dropout(0.5)(x)
x = BatchNormalization()(x)

output = Dense(self.num_classes, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=output)
```


RESNET-50

ACTUAL: Hyundai Tucson SUV 2012



ACTUAL: Land Rover LR2 SUV 2012



ACTUAL: Audi 100 Sedan 1994



PREDICTION: Jeep Liberty SUV 2012



PREDICTION: Ford Freestar Minivan 2007



PREDICTION: Volvo 240 Sedan 1993

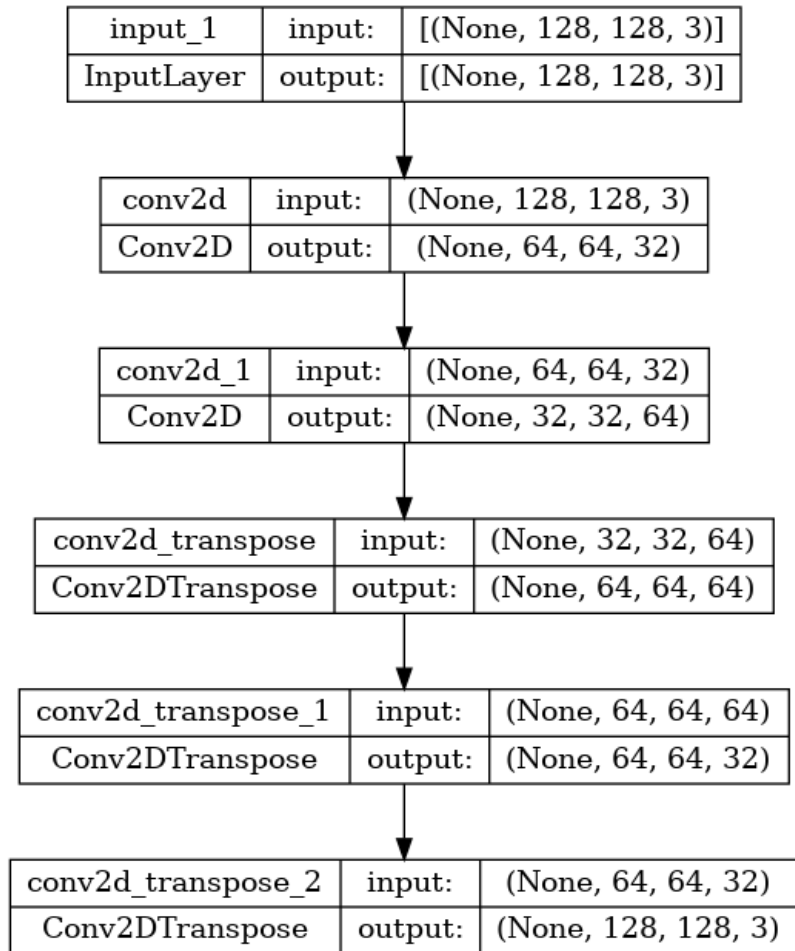


Autoencoder

Perturbación de las imágenes con ruido



ARQUITECTURA



loss: 6.9119e-04 - accuracy: 0.7632 - val_loss:
6.0662e-04 - val_accuracy: 0.7647

RESULTADOS

